



# Recent Advances in the OSQP Solver

Ian McInerney – Imperial College London  
Bartolomeo Stellato – Princeton University, ORFE  
Vineet Bansal – Princeton University, CSML  
Amit Solomon – Princeton University

**2025 INFORMS Annual Meeting**  
**28/10/2025**

# Main contributors

Ian McInerney



Vineet Bansal



Bartolomeo Stellato



Paul Goulart



Goran Banjac



Michel Schubiger



# Agenda

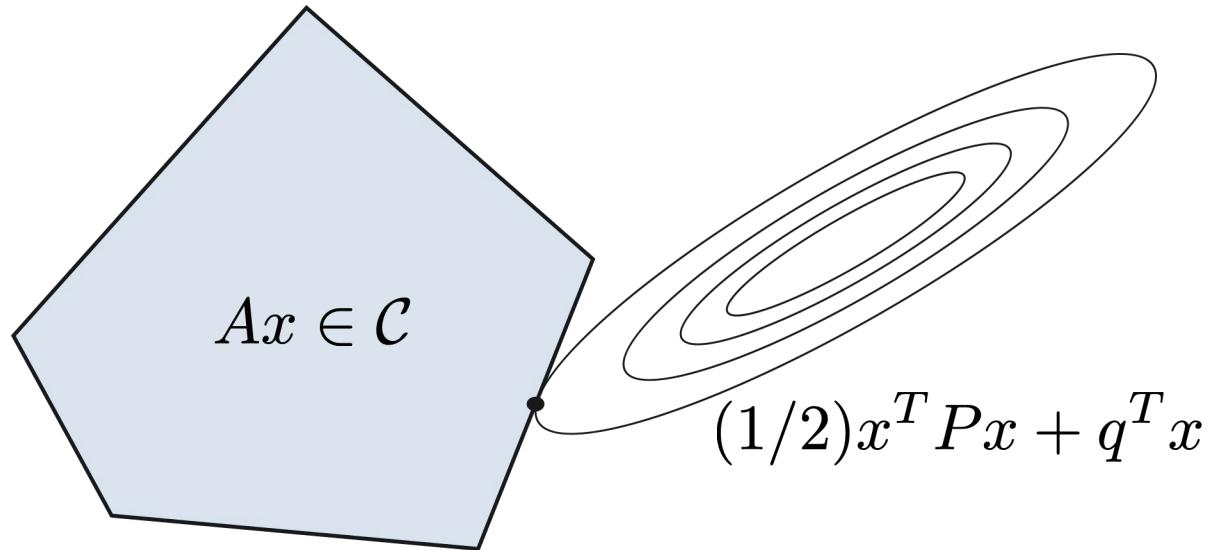
- The OSQP Solver
- New features:
  - Duality gap termination criteria
  - cuDSS integration
- The future

# The OSQP Solver

# Quadratic Programming Problem

$$\begin{aligned} & \text{minimize} && (1/2)x^T Px + q^T x \\ & \text{subject to} && Ax \in \mathcal{C} \end{aligned}$$

Quadratic program:  $\mathcal{C} = [l, u]$



# The OSQP Algorithm

## Problem

$$\begin{aligned} & \text{minimize} && (1/2)x^T Px + q^T x \\ & \text{subject to} && l \leq Ax \leq u \end{aligned}$$

## Algorithm

**Linear system  
solve**

$$(x^{k+1}, \nu^{k+1}) \leftarrow \text{solve} \begin{bmatrix} P + \sigma I & A^T \\ A & -\frac{1}{\rho} I \end{bmatrix} \begin{bmatrix} x^{k+1} \\ \nu^{k+1} \end{bmatrix} = \begin{bmatrix} \sigma x^k - q \\ z^k - \frac{1}{\rho} y^k \end{bmatrix}$$

$$\tilde{z}^{k+1} \leftarrow z^k + (\nu^{k+1} - y^k)/\rho$$

$$z^{k+1} \leftarrow \Pi(\tilde{z}^{k+1} + y^k/\rho)$$

$$y^{k+1} \leftarrow y^k + \rho(\tilde{z}^{k+1} - z^{k+1})$$

# OSQP Solver Features

-  Infeasibility detection
-  Solution Polishing
-  Matrix & vector updates
-  Modular linear algebra (CUDA & MKL)
-  Embedded (library-free) mode
-  Solver instance code generation
-  High-level interfaces (JuMP, CVXPY, MATLAB)
-  Open source – Apache 2.0 license

# Using MKL & CUDA from Python

## One-line import change

```
# Import OSQP from a specific algebra backend module
from osqp.mkl import OSQP as OSQP_mkl
from osqp.cuda import OSQP as OSQP_cuda

prob_mkl = OSQP_mkl()
prob_cuda = OSQP_cuda()

# Setup workspace and change alpha parameter
prob_mkl.setup(P, q, A, l, u, alpha=1.0)

# Solve problem
res = prob_mkl.solve()
```

## Setting in object constructor

```
# Create an OSQP object with a specific algebra backend
if osqp.algebra_available('cuda'):
    # 'builtin' (default), 'mkl', or 'cuda'
    prob = osqp.OSQP(algebra='cuda')
else:
    prob = osqp.OSQP()
```

```
# Setup workspace and change alpha parameter
prob.setup(P, q, A, l, u, alpha=1.0)

# Solve problem
res = prob.solve()

...
```

```
# Solve with OSQP cuda on CVXPY
import cvxpy as cp

problem = cp.Problem(...)
problem.solve(solver=OSQP, algebra="cuda")
```

It works  
with CVXPY



# New features

Duality Gap Termination

# Three important residuals

## Primal Residual

Measure of primal feasibility

$$r_{\text{prim}} = Ax - z$$

## Dual Residual

Measure of dual feasibility

$$r_{\text{dual}} = Px + q + A^T y$$

## Duality Gap

Distance between primal and dual objectives

$$r_{\text{gap}} = x^T Px + q^T x + u^T y_+ + l^T y_-$$

# Duality gap termination - Accuracy

**Residual-only termination failed to minimize duality gap in 27.8% of benchmark problems**

Equality-constrained QP	3.0%
Huber	73.5%
Lasso	0.0%
Portfolio Optimization	100%
Random QP	9%
MPC Control Prob.	5.5%
SVM	0.0%

Percent of problems failing to reduce duality gap to below 1e-3 tolerance

# Duality gap termination

- Implemented a duality gap termination criteria
- Iterate until duality gap is less than the tolerance:

$$|r_{\text{gap}}| \leq \epsilon_{\text{gap}} = \epsilon_{\text{abs}} + \epsilon_{\text{rel}} \max\{|x^T Px|, \|q^T x\|, |S_C(y)|\}.$$

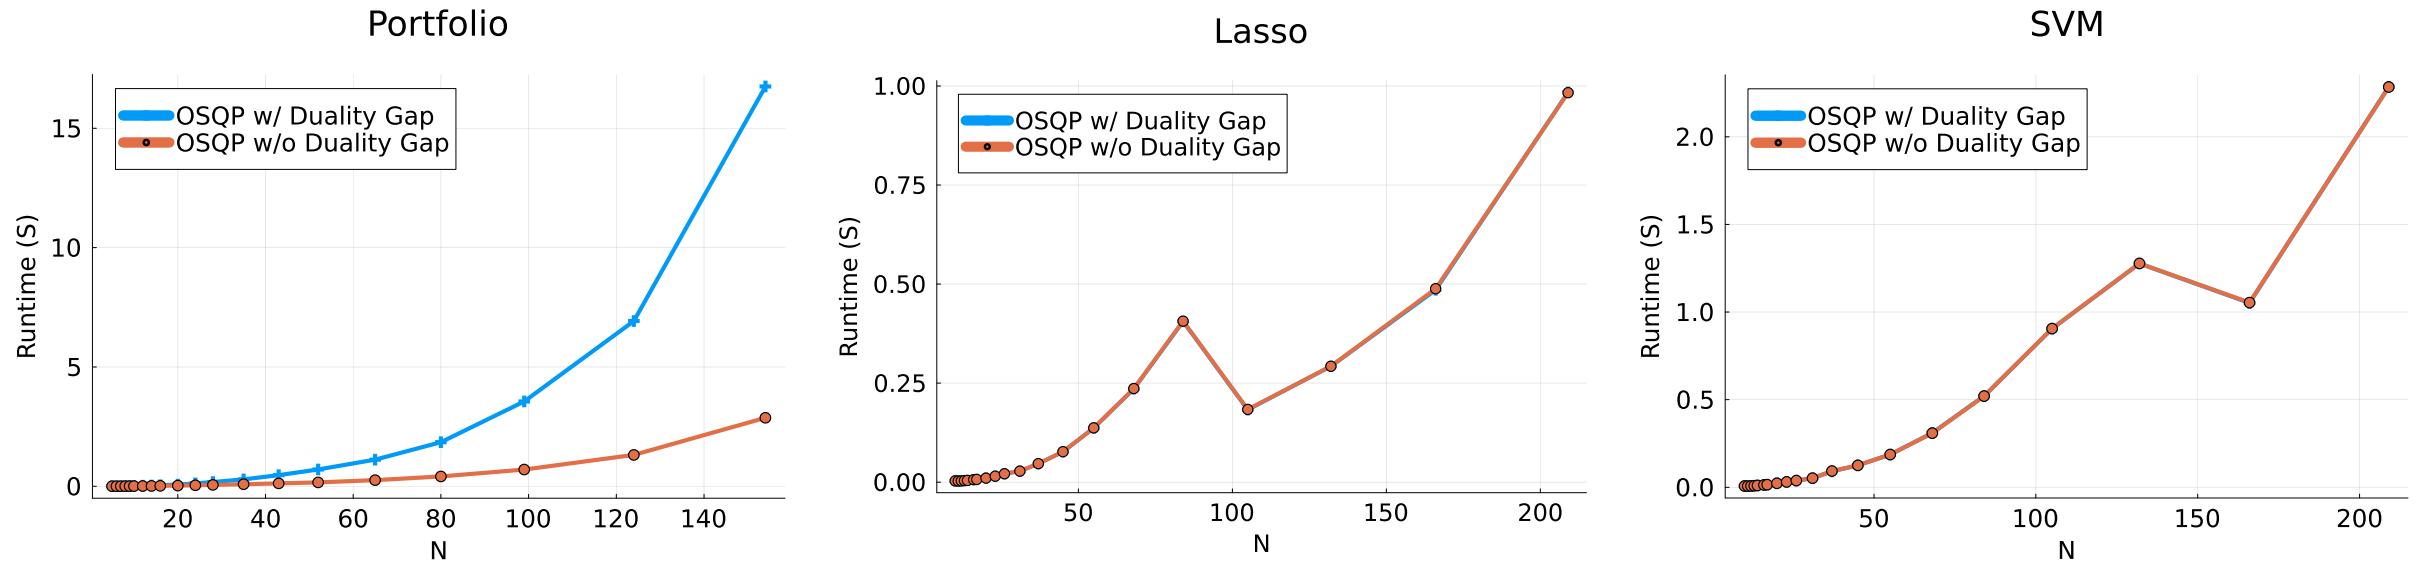
# Duality gap termination - Accuracy

With all 3 residuals checked, all benchmark problems converged to below the duality gap tolerance

	<b>OSQP &lt;1.0</b>	<b>OSQP 1.0</b>
Equality-constrained QP	3.0%	0.0%
Huber	73.5%	0.0%
Lasso	0.0%	0.0%
Portfolio Optimization	100%	0.0%
Random QP	9%	0.0%
MPC Control Prob.	5.5%	0.0%
SVM	0.0%	0.0%

Percent of problems failing to reduce duality gap to below 1e-3 tolerance

# Duality gap termination - Performance



Solved to low accuracy: 1e-3

CPU: Intel Xeon W-2255 @ 3.7GHz

# New features

## cuDSS Integration

# Solving the linear system

## Direct method

**Quasi-definite  
matrix**

$$\begin{bmatrix} P + \sigma I & A^T \\ A & -\frac{1}{\rho} I \end{bmatrix} \begin{bmatrix} x \\ \nu \end{bmatrix} = \begin{bmatrix} \sigma x^k - q \\ z^k - \frac{1}{\rho} y^k \end{bmatrix}$$

Solve using  
LDL  
factorization

## Indirect method

**Positive-definite  
matrix**

$$(P + \sigma I + \rho A^T A) x = \sigma x^k - q + A^T (\rho z^k - y^k)$$

Solve using  
conjugate  
gradient

# CUDA Direct Sparse Solver (cuDSS) Integration

## LDL solver from cuDSS

Solution of 2x2 KKT system:

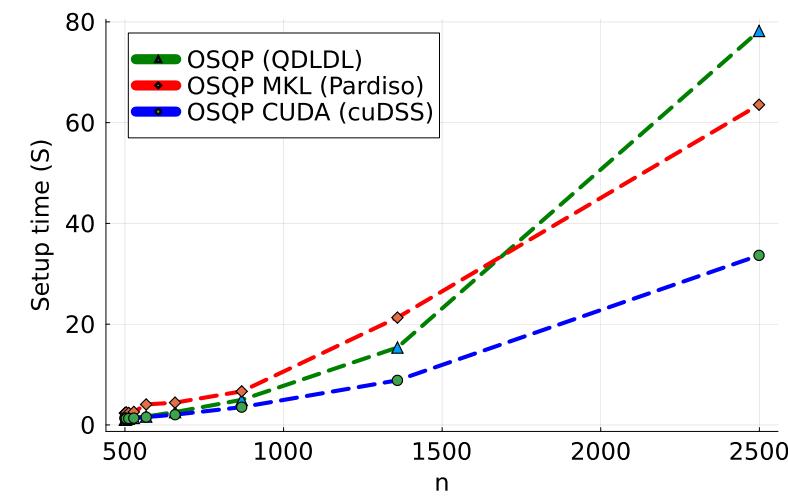
$$\begin{bmatrix} P + \sigma I & A^T \\ A & -\frac{1}{\rho} I \end{bmatrix} \begin{bmatrix} x \\ \nu \end{bmatrix} = \begin{bmatrix} \sigma x^k - q \\ z^k - \frac{1}{\rho} y^k \end{bmatrix}$$

## GPU-native solver

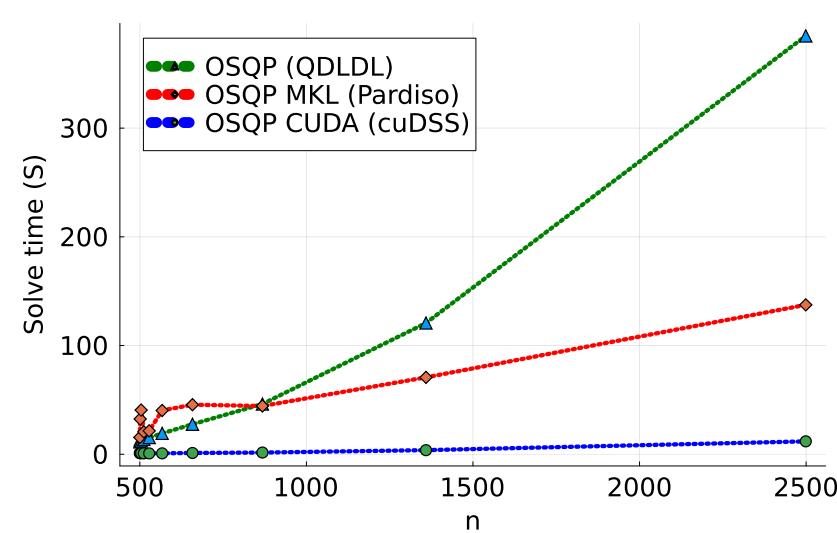
- Form 2x2 KKT matrix structure on CPU in setup, then transfer to GPU
- All solve computations are on GPU

# cuDSS – Performance vs. other direct solvers

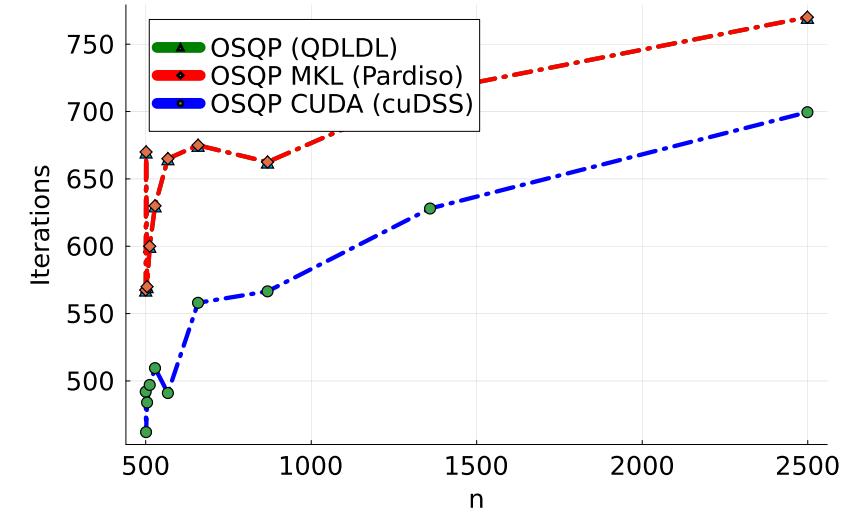
SVM



SVM



SVM



# of variables in problem =  $101 \cdot n$   
# of constraints in problem =  $200 \cdot n$   
Size of KKT matrix =  $301 \cdot n \times 301 \cdot n$

Solved to low accuracy: 1e-3

# Linear System Solvers Available

	Built-in	MKL	CUDA
Direct	(QDLDL)	(Pardiso)	(cuDSS)
Indirect		(PCG)	(PCG)

# The future

# Future improvements

## Software Engineering

- CUDA Streams
- CUDA Graphs
- Batched mode

## Algorithmic Improvements

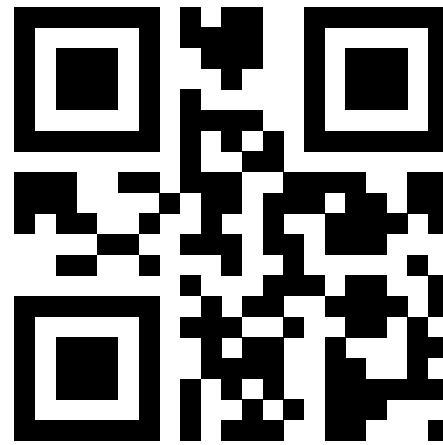
- Restarting
- Mixed/low precision
- ADMM-tailored solution derivatives
- Improved step-size selection strategies

# Learn more & get OSQP

## C Library



[github.com/osqp/osqp](https://github.com/osqp/osqp)



[osqp.org](https://osqp.org)

## Python

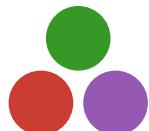


`pip install osqp[mkl,cu12]`



[osqp.org/docs/](https://osqp.org/docs/)

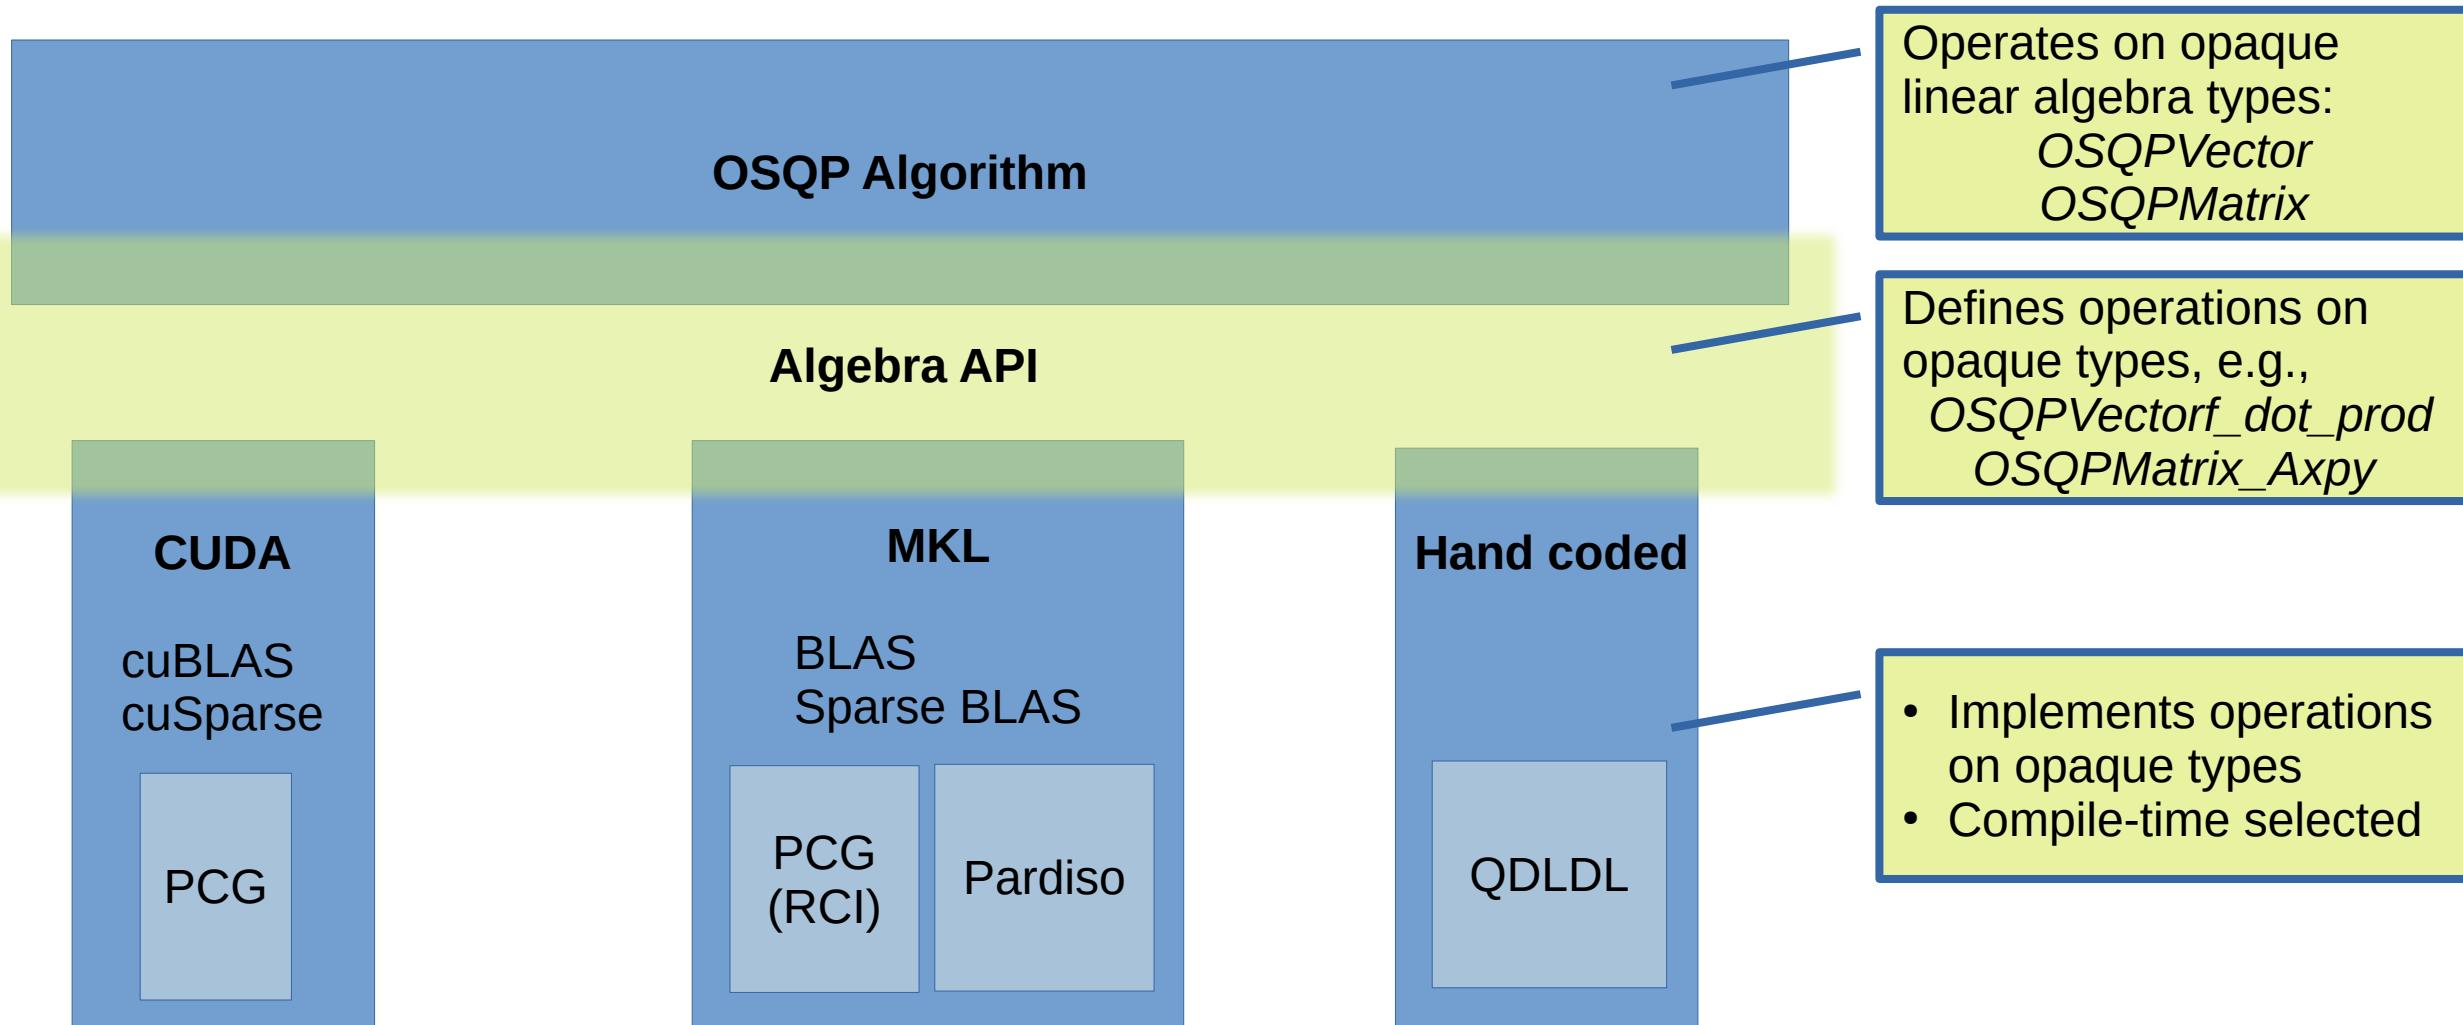
## Julia



`Pkg.add(["OSQP", "OSQPMKL", "OSQPCUDA"])`

# Backup slides

# Modular Linear Algebra



# Modular Linear Algebra – Key Points

## Builtin

- Hand coded
- CSC-based sparse matrices
- Library free
- Linear system solvers
  - QDLDL

## MKL

- CSC-based sparse matrices
- BLAS vector operations
- Linear system solvers
  - Pardiso
  - RCI Preconditioned conjugate gradient

## CUDA

- cuSparse & cuBLAS libraries
- CSC & CSR matrices
- All data fully GPU resident
  - *osqp\_setup* – Data copied to OSQP GPU workspace
  - *osqp\_solve* – CPU-managed control flow, only transfer status values
- Linear system solvers
  - Preconditioned conjugate gradient



Same OSQP API for all backends

# Modular Linear Algebra from Julia

## One-line import change

```
● ● ●

using JuMP
using OSQP
using OSQPMKL

model = Model( () -> OSQP.Optimizer(OSQPMKLAlgebra()) )

@variable(model, x >= 0)
@variable(model, 0 <= y <= 3)
@objective(model, Min, 12x + 20y)
@constraint(model, c1, 6x + 8y >= 100)
@constraint(model, c2, 7x + 12y >= 120)
print(model)
optimize!(model)
```

It works  
with JuMP